

HTTP


Daniel Zappala

CS 360 Internet Programming
Brigham Young University


HTTP Objects

Requesting Objects

- clients request **objects** from servers using the HTTP protocol
 - client sends an HTTP request
 - server sends an HTTP response
- does not necessarily have a GUI
 - collecting hourly reports on competitor's prices
 - mining Twitter data



Web Objects



- **object names:** Uniform Resource Identifier (URI)
 - a name that refers to a resource
 - a Uniform Resource Locator (URL) is one type of URI
 - popular URL schemes: http, ftp, gopher, mailto
- **object delivery:** Hypertext Transfer Protocol (HTTP)
 - IETF standard
 - defines message format for making requests and receiving responses
- **object format:** Hypertext Markup Language (HTML)
 - representation of documents in ASCII format
 - many other formats - XHTML, XML, PNG, JPG, PDF, etc.

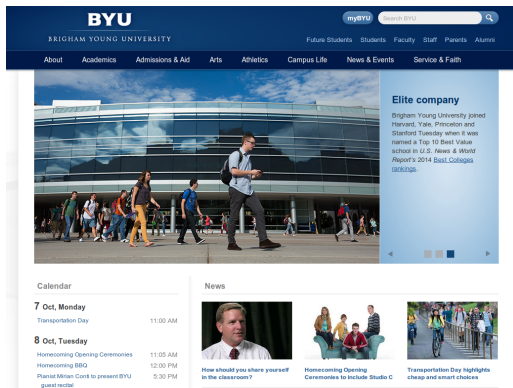
URIs, URLs, and URNs



- URI: Uniform Resource Identifier
 - The generic set of all names that refer to resources
- URL: Uniform Resource Locator
 - An informal term (no longer used in technical specifications) associated with popular URI schemes: http, ftp, gopher, mailto, etc.
- URN : Uniform Resource Name
 - A URI that has an institutional commitment to persistence, availability, etc. May also be a URL: see PURLs.
 - persistent, location-independent resource identifiers, urn: specified by RFC 2141

Container Objects

- a web page consists of a **container** object, which may link to other objects
- fetching a web page consists of requesting the container object and then requesting any linked objects (images, CSS, Javascript, etc.)



Handling Objects

- determines how responses are handled
 - appearance (fonts)
 - content transformations (language)
 - whether to accept cookies
 - whether to allow javascript, popups
 - MIME types and handlers
- see Chrome advanced settings

HTTP Message Formats

HTTP Standards

- HTTP 1.0
 - RFC 1945: <http://www.ietf.org/rfc/rfc1945.txt>
 - very basic protocol, documenting what earliest servers and browsers used
- HTTP 1.1
 - RFC 2616: <http://www.ietf.org/rfc/rfc2616.txt>
 - backward compatibility with HTTP/1.0, plus many improvements and features
 - what all modern servers and browsers uses

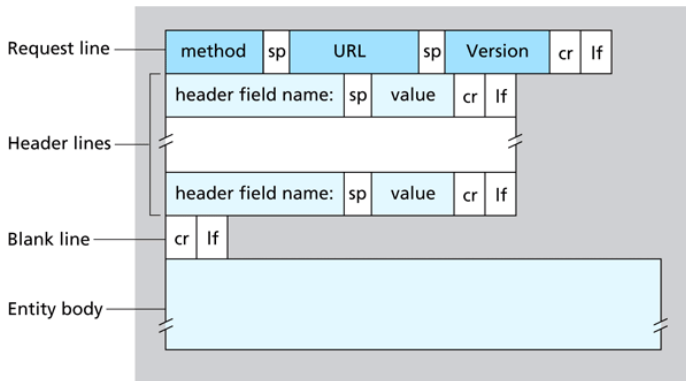
Standardization History

- HTTP/1.1 standardization took 4 years
- after two years RFC 2068 captured the early state of the HTTP/1.1 specification *process*
- browsers started implementing the “HTTP/1.1 standard” before it was officially a standard
- HTTP/1.1 needed to be backward-compatible with many browsers or else many sites would not deploy it
- as a result, RFC 2616 contains some idiosyncrasies, never fully adopted as an official standard

Specification Language

- specification language is precise
 - **MUST**: absolutely essential - if you don't implement this feature you are not compliant
 - **SHOULD**: recommendation - you are compliant if you don't implement this feature, but you should implement it if at all possible
 - **MAY**: optional - not considered necessary
 - there are two obvious counterparts: **MUST NOT**, **SHOULD NOT**
- see RFC 2119

HTTP Request Format

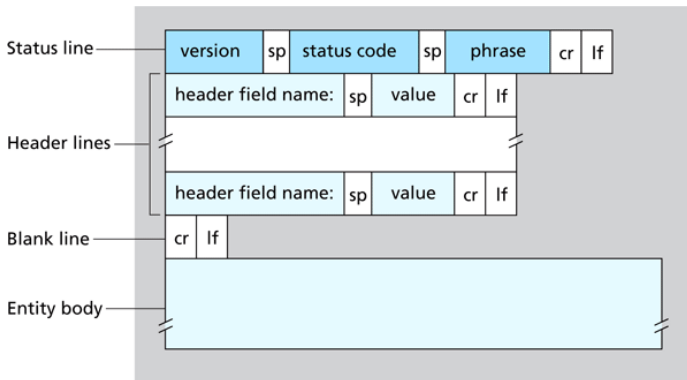


- entity body is optional, with a header that indicates the length of the body in bytes

Example HTTP Request

```
1 GET /index.html HTTP/1.1
2 Host: cs360.byu.edu
3 Accept: text/html,application/xhtml+xml,application/xml
4 Accept-Encoding: gzip, deflate, sdch
5 Accept-Language: en-US,en;q=0.8
6 Connection: keep-alive
7 User-Agent: Mozilla/5.0 (X11; Linux x86_64) AppleWebKit
8 /537.36 (KHTML, like Gecko) Chrome/29.0.1547.76 Safari/537.36
```

HTTP Response Format



- entity body is optional, with a header that indicates the length of the body in bytes

Example HTTP Response

```
1 HTTP/1.1 200 OK
2 Connection: Keep-Alive
3 Date: Mon, 07 Oct 2013 16:31:00 GMT
4 ETag: "c01e0-c4f-4e6c0444f4e40"
5 Keep-Alive: timeout=5, max=100
6 Server: Apache/2.2.22 (Ubuntu)
7 Vary: Accept-Encoding
8
9 <!DOCTYPE html>
10 <html lang="en">
11   <head>
12     <title>CS 360 Internet Programming</title>
13   ...
```

- use [telnet cs360.byu.edu 80](http://telnet.cs360.byu.edu) to experiment

Request Methods

GET, HEAD

- GET
 - see RFC 1945, section 8.1
 - MAY encode arguments to a script
 - `GET /book.cgi?lastname=Stevens`
- HEAD
 - see RFC 1945, section 8.2
 - MUST NOT return entity body

POST

- see RFC 1945, section 8.3
- MUST have a [Content-Length](#) to determine length of entity body
- actions taken depend on the URI

PUT, DELETE, LINK, UNLINK

- see RFC 1945, Appendix D
- PUT and DELETE considered dangerous because the client is in charge
- not usually supported by web servers
- when coupled with authentication, can create an API for third-party clients to interact with web services

Headers

Header Syntax

1 HTTP-header = field-name ":" [field-value] CRLF

- arbitrary length
- order is not significant, but recommended that you send General, Request, Response, and Entity Headers in that order
- see RFC 1945, section 4.2

General Headers

- **Date**
 - RFC 1945, Section 10.6
 - date and time at which message originated
 - specifies format of date and time
 - MUST be in GMT

Request Headers

- **If-Modified-Since**
 - RFC 1945, Section 10.9
 - defines a conditional GET - retrieve resource only if it hasn't been changed since the given date
 - uses the same format as the **Date** header
- **From**
 - RFC 1945, Section 10.8
 - used to give user's email address
 - server should parse and log
- **User-Agent**
 - RFC 1945, Section 10.15
 - information about user's browser
 - server should parse and log

Response Headers

- **Server**
 - RFC 1945, Section 10.14
 - identifies server software
 - many servers include this in all responses

Entity Headers

- **Content-Length**
 - RFC 1945, Section 10.4
 - size of the entity body in bytes
 - MUST be included for any message containing an entity body
 - server should send an error if it is missing, see Section 7.2.2
- **Last-Modified**
 - RFC 1945, Section 10.10
 - date and time at which resource was last modified
 - MUST NOT send a time later than message origination
- **Content-Type**
 - RFC 1945, Section 10.5
 - indicates media type of entity body
 - server SHOULD include this header if the message has an entity body

Response Codes

Response Classes

- HTTP groups response codes together into classes
- 100: informational class
- 200: success class
- 300: redirection class
- 400: client error class
- 500: server error class

Notable Error Codes

- 200 OK
- 304 Not Modified
- 400 Bad Request
- 403 Forbidden
- 404 Not Found
- 501 Not Implemented

Features

Multiple Web Sites

- HTTP/1.1 allows many web sites per server
- server MUST include `Host` header in HTTP/1.1

1	GET /index.html HTTP/1.1
2	Host: cs360.byu.edu

Content Negotiation

- client can ask for different versions
- server-driven
 - client sends hints about user's preference using `Accept-Language`, `Accept-Charset`, `Accept-Encoding`
 - server chooses the best match
- agent-driven
 - server responds with `300 Multiple Choices` that includes list of available representations
 - client makes a new request with the chosen variant
- server-driven is widely used, reduces latency

Persistent Connections and Pipelining

- TCP works best for long-lived connections
 - starts with a slow rate
 - increases rate as segments are delivered successfully
 - decreases rate when loss occurs
- HTTP allows browser to use a single connection for many requests, to keep the rate high
 - useful for fetching images, CSS, Javascript
 - can also use to fetch subsequent pages and their objects
- Connection header
 - by default, connections stay open
 - server sends `Connection:close` to notify client it will close the connection
 - connections may be closed at any time
- many clients use parallel connections to further increase speed
- Chrome uses prefetching to guess which page you will view next

Downloading a Part of a File

- user may want only part of a resource
 - continuing after an aborted transfer
 - fetching in parallel from multiple servers
- **Range** header
 - specify one or more ranges of contiguous bytes
 - **Range:** `bytes=0-1000`
 - if server's response contains a range, it uses **206 Partial Content**

Compression

- 1997 study showed that compression could save 40% of the bytes sent via HTTP
- **Content-Encoding**: indicates end-to-end content coding
- **Transfer-Encoding**: indicates hop-by-hop transfer-codings (applied by proxies)
- **Accept-Encoding**: indicates the type of encodings a client can accept
- **TE**: indicates the type of transfer-codings a client prefers
- IANA registers transfer-coding token values, including **chunked**, **identity**, **gzip**, **compress**, and **deflate**

Chunked Transfers

- chunked transfers are used by a web server for dynamic content, when it doesn't know ahead of time the length of the entire body of the object
- does **not** use Content-Length header
- sends the body as a series of chunks, each having the following format:
 - length of the chunk, in hexadecimal, followed by a CRLF (`\r\n`)
 - a sequence of bytes, whose length is given by the previous line
 - a CRLF
- to end the body, the server sends a valid last chunk with a length of 0 bytes

Cookies

Motivation

- HTTP does not keep track of state, but servers need it
- web sites would like to
 - greet you personally
 - provide an account
 - track your shopping cart
 - provide product recommendations
 - store state about your email account
 - invade your privacy – track what you like to search for, what you like to buy, what you like to read, etc.

How a Cookie Works: RFC 2109

- ① user sends usual HTTP request message
- ② server creates a unique ID for this user
- ③ server returns usual HTTP response, but with a cookie header
 - `Set-Cookie: session-id=1678`
- ④ when user sends subsequent HTTP requests, they also include a cookie header
 - `Cookie: session-id=1678`
- ⑤ server can customize its response based on cookie value

Security Implications


- session hijacking
 - 1 user supplies name and password
 - 2 server sets a cookie, which user passes to server for each transaction
 - 3 intruder snoops on your cookie (*ethereal* on a shared LAN) or guesses its value
 - 4 intruder supplies cookie to web server and learns your state (e.g. a credit card)
- precautions
 - use SSL to authenticate the user each time
 - send the cookie over the encrypted connection so it can't be snooped
 - use a cookie that expires once the session is finished so it can't be replayed

SPDY, HTTP/2, and QUIC

SPDY

- developed by Google (if in Chrome, will work with any Google web server)
- **multiplexing**: streams of requests and responses can be interleaved on a single connection (HTTP requires sequential/pipelined requests)
- **header compression**: avoid repeated headers in multiple requests/responses
- **request prioritization**: client can mark high priority requests
- **server hint**: can indicate high priority resources
- **server push**: can push additional resources (e.g. CSS, JavaScript) in response to a request (e.g. index.html)

SPDY Multiplexing



- for speedup results:
<https://www.chromium.org/spdy/spdy-whitepaper>

HTTP/2


- IETF adopted SPDY improvements
- server push requires a subscription from the client, rather than unilateral action by the server
- binary framing

HTTP/2 Binary Framing

HTTP/2 Binary Framing

Enabled by dumping newline delimited ASCII


- Frame Header



- Types

- HEADERS, DATA, PRIORITY, RST_STREAM, SETTINGS, PUSH_PROMISE, PING, GOAWAY, WINDOW_UPDATE, CONTINUATION

QUIC



QUIC

- low latency connection establishment
 - cache TLS credentials to resume TLS sessions
- more rapidly change transport layer (implemented in application)
 - loss recovery
 - congestion control
 - FEC
- connection migration